

Meridian & Associates LLP

Global Professional Services

AI Governance & Responsible Technology Policy

Version 1.0

Effective Date: February 2026

Classification: Confidential

Prepared by the Office of Strategic Communications

Meridian & Associates LLP

CONFIDENTIAL

This document contains proprietary information. Do not distribute outside of authorized RFP response teams without written consent.

Table of Contents

1	AI Ethics Framework and Principles
2	Generative AI Acceptable Use Policy
3	Client Data Protection in AI Systems
4	Model Risk Management Framework
5	Algorithmic Bias Monitoring and Fairness
6	AI Output Review Requirements
7	Third-Party AI Tool Vetting
8	AI-Specific Incident Response
9	Governance Structure
10	Training Requirements

1. AI Ethics Framework and Principles

Meridian & Associates LLP recognizes that artificial intelligence and machine learning technologies present both transformative opportunities and significant risks. This policy establishes the firm's framework for the responsible development, deployment, and use of AI systems across all service lines, internal operations, and client engagements. It applies to all partners, employees, contractors, and third parties acting on behalf of Meridian.

Meridian's Responsible AI Principles guide all AI-related activities:

- **Beneficence** -- AI systems must be designed and deployed to create genuine value for clients, employees, and society. AI should augment human capabilities, not replace human judgment in contexts requiring ethical reasoning, empathy, or professional skepticism.
- **Transparency** -- Clients, regulators, and affected individuals must be able to understand how AI systems influence decisions that affect them. Meridian discloses AI use in client deliverables where material.
- **Fairness** -- AI systems must not discriminate against individuals or groups based on protected characteristics. Outputs must be regularly tested for disparate impact.
- **Accountability** -- Every AI system deployed at Meridian must have a designated human owner who is accountable for its outputs, performance, and compliance.
- **Privacy** -- AI systems must respect data protection rights. Personal data is never used for model training without explicit consent and lawful basis.
- **Safety and Reliability** -- AI systems must be tested rigorously before deployment. Failure modes must be understood, documented, and mitigated.
- **Security** -- AI systems must be protected against adversarial attacks, data poisoning, prompt injection, and model theft.

2. Generative AI Acceptable Use Policy

The following guidelines govern the use of generative AI (GenAI) tools -- including large language models, image generators, and code generation systems -- by all Meridian professionals:

2.1 Permitted Uses

- Internal research, brainstorming, and ideation support using approved tools only.
- First-draft generation for internal documents, proposals, and communications (subject to human review before distribution).
- Code development assistance (code completion, refactoring suggestions, test generation) using approved IDE-integrated tools.
- Data analysis acceleration (summarization, pattern identification, visualization) on non-confidential or appropriately classified data.
- Translation and language assistance for internal and client communications.
- Knowledge retrieval from Meridian's proprietary knowledge bases via MeridianAI.

2.2 Prohibited Uses

- Inputting client confidential, restricted, or personal data into any AI tool not on the Approved AI Tools Registry -- including personal accounts for ChatGPT, Gemini, Claude, Copilot, or any other public GenAI service.

- Using AI-generated outputs as final work product delivered to clients without qualified human review and validation.
- Relying on AI for audit opinions, legal advice, tax computations, regulatory filings, or any deliverable requiring professional certification or sign-off.
- Using AI to generate content intended to deceive, mislead, or fabricate information.
- Using AI to circumvent security controls, access restrictions, or compliance policies.
- Using AI to make or materially influence employment, promotion, or compensation decisions without meaningful human oversight.
- Using AI-generated code in production environments without security review and testing.

3. Client Data Protection in AI Systems

Protecting client data is paramount in all AI deployments. The following controls are mandatory:

- No client data in training: Client data is never used to train, fine-tune, or improve any AI/ML model -- whether proprietary or third-party. This is an absolute prohibition enforced through technical controls (data classification enforcement, training pipeline isolation) and contractual terms with all AI vendors.
- Opt-in policy: Clients must provide explicit written consent before their engagement data is processed by any AI system beyond basic search and retrieval. Consent is documented in the engagement letter or a supplemental AI processing addendum.
- Data minimization: Only the minimum data necessary for the specific AI task is provided to the system. PII is redacted or pseudonymized wherever technically feasible.
- Ephemeral processing: Client data processed by AI systems for inference is not persisted beyond the active session. No client data is retained in AI system logs, caches, or vector stores beyond 24 hours.
- Data residency: AI processing of client data respects all data residency requirements specified in the engagement agreement. Regional AI deployment options available for EU, UK, APAC, and U.S.-only processing.
- Vendor contractual controls: All AI vendor contracts include explicit prohibitions on using client data for model training, improvement, or any purpose beyond the contracted inference service. Verified annually through vendor audit or SOC 2 review.

4. Model Risk Management Framework

Meridian applies a risk-based approach to AI/ML model governance, aligned with the principles of SR 11-7 (OCC/Fed model risk management guidance) and adapted for professional services context:

4.1 Model Classification

- Tier 1 -- High Risk: Models whose outputs directly influence client deliverables, financial decisions, regulatory outcomes, or hiring/promotion decisions. Subject to full validation cycle, independent review, and annual re-validation.
- Tier 2 -- Medium Risk: Models used for internal analytics, resource planning, or quality assurance. Subject to initial validation and biennial review.
- Tier 3 -- Low Risk: Models used for productivity assistance (summarization, translation, code completion). Subject to initial assessment and periodic spot checks.

4.2 Model Lifecycle Controls

- Model Registry: All AI/ML models in use across the firm are cataloged in a central Model Registry with metadata including: purpose, owner, data inputs, known limitations, risk tier, validation status, and deployment environment.
- Development standards: All proprietary models developed using Meridian's ML Engineering Standards, which mandate version control, reproducibility, documentation, and peer review.
- Validation: Independent validation by the AI Risk team (separate from the model development team) for all Tier 1 and Tier 2 models. Validation covers: conceptual soundness, data quality, performance metrics, limitations analysis, and ongoing monitoring plan.
- Monitoring: Production models monitored for data drift, performance degradation, and anomalous outputs. Automated alerts trigger review when metrics exceed defined thresholds.
- Retirement: Documented decommissioning process including stakeholder notification, archival of model artifacts, and transition plan for dependent processes.

5. Algorithmic Bias Monitoring and Fairness

Meridian is committed to ensuring that AI systems do not perpetuate or amplify unfair bias. The following requirements apply to all AI systems that produce outputs affecting individuals:

- Pre-deployment testing: All Tier 1 models undergo fairness testing before deployment, analyzing outputs across protected characteristics (race, gender, age, disability, etc.) where applicable. Fairness metrics include demographic parity, equalized odds, and calibration across subgroups.
- Ongoing monitoring: Production models with human-impacting outputs are monitored quarterly for disparate impact. Results reported to the AI Ethics Board.
- Training data audits: Data used to train or fine-tune proprietary models is audited for representational balance and potential sources of historical bias.
- Remediation: Models found to exhibit statistically significant bias are flagged for immediate review. Remediation options include: retraining with balanced data, algorithmic debiasing techniques, output adjustment, or model retirement.
- Client transparency: For client engagements where AI-assisted analytics or decision-support systems are delivered, fairness testing methodology and results are included in the deliverable documentation upon client request.
- External audit: Annual independent audit of the firm's highest-risk AI systems by a qualified external assessor. Results reported to the Board of Partners.

6. AI Output Review Requirements

Meridian mandates human-in-the-loop review for all AI-generated or AI-assisted outputs that are delivered to clients or used in consequential decision-making:

- Client deliverables: All AI-generated or AI-assisted content included in client deliverables must be reviewed and validated by a qualified professional before delivery. The reviewer must be competent in the subject matter and must attest that the output is accurate, appropriate, and free from hallucination or material error.
- Audit and assurance: AI tools may assist with data analysis, anomaly detection, and sampling in audit engagements, but all audit conclusions, opinions, and findings must be determined by qualified auditors exercising professional judgment. AI outputs in audit are treated as evidence to be evaluated, not as conclusions.
- Legal and regulatory: AI-generated legal analysis, regulatory interpretations, or tax computations must be reviewed

by a licensed professional before use.

- Code and technology: AI-generated code must pass the same code review, security scanning (SAST/DAST), and testing standards as human-written code before production deployment.
- Internal decisions: AI-assisted recommendations for hiring, promotion, performance evaluation, or resource allocation must be reviewed by the responsible human decision-maker, who retains full authority to accept, modify, or reject the AI recommendation.
- Documentation: The use of AI in producing any deliverable must be documented in the engagement file, including the tools used, the nature of AI assistance, and the reviewer who validated the output.

7. Third-Party AI Tool Vetting

All AI/ML tools and services from external vendors must undergo a formal assessment before approval for use within Meridian:

- Assessment process: Requests for new AI tools are submitted to the AI Governance Board via the Technology Procurement Portal. Assessment includes: security review (InfoSec team), privacy review (Privacy Office), legal review (General Counsel), and ethical risk assessment (AI Ethics Board).
- Security requirements: SOC 2 Type II certification (or equivalent). Penetration test results within the last 12 months. Vulnerability disclosure policy. No training on customer data without explicit opt-in.
- Privacy requirements: GDPR/CCPA-compliant data processing terms. Clear disclosure of data flows, sub-processors, and retention policies. Data residency options compatible with Meridian's client requirements.
- Contractual requirements: Explicit prohibition on using Meridian or client data for model training. Right to audit. Breach notification within 24 hours. Data deletion upon termination.
- Ongoing monitoring: Approved tools are reassessed annually. Material changes to vendor terms of service, data practices, or security posture trigger interim review.
- Approved AI Tools Registry: Maintained on the firm intranet with current list of 34 approved AI tools and services, categorized by use case and data classification level.

8. AI-Specific Incident Response

AI-related incidents require specialized response procedures in addition to the firm's general Incident Response framework:

- Scope: AI-specific incidents include: AI-generated output that is materially incorrect and has been delivered to a client; discovery of bias in a deployed model; unauthorized data exposure through an AI system; prompt injection or adversarial manipulation; unintended disclosure of confidential information via AI-generated output.
- Reporting: AI incidents must be reported to both the IT Service Desk (for general tracking) and the AI Ethics Board (ai-ethics@meridianllp.com) within 2 hours of discovery.
- Containment: For AI output errors affecting client deliverables, the engagement partner is notified immediately. Affected outputs are recalled or corrected. Client notification per engagement terms.
- Investigation: Root cause analysis conducted by the AI Risk team in collaboration with the development/deployment team. Focus areas: data quality, model behavior, prompt engineering, human oversight adequacy.
- Remediation: Corrective actions may include: model retraining, guardrail adjustment, additional human review requirements, tool restriction, or vendor escalation.

- Tracking: All AI incidents logged in the firm's incident management system with dedicated AI incident category. Quarterly trend reports provided to the AI Ethics Board and CISO.
- FY2025 summary: 28 AI-related incidents reported; 23 classified as low severity (output quality); 5 classified as medium (data handling); zero classified as high; average resolution time: 3.2 business days.

9. Governance Structure

Meridian's AI governance is led by dedicated leadership and oversight bodies:

9.1 Chief AI Officer

Dr. Amara Osei serves as Meridian's Chief AI Officer (a dual role with Chief Data Officer), reporting to the CTO with a dotted line to the Global Managing Partner. The Chief AI Officer is responsible for AI strategy, governance policy, risk oversight, and the firm's AI research and innovation agenda.

9.2 AI Ethics Board

- Composition: 12 members including the Chief AI Officer (Chair), CISO, General Counsel, Chief Risk Officer, DPO, representatives from each service line, and two external advisors (a computer science professor and a civil liberties attorney).
- Mandate: Approve the AI Governance Policy and material updates; review and approve Tier 1 AI use cases; adjudicate escalated ethical concerns; oversee fairness testing program; commission external AI audits.
- Meeting cadence: Monthly, with ad hoc sessions for urgent matters.
- Authority: The AI Ethics Board has the authority to suspend or prohibit any AI use case, tool, or vendor that does not comply with this policy, regardless of commercial impact.

9.3 AI Risk Team

- 18 dedicated professionals within Risk & Compliance responsible for: model validation, bias testing, AI incident investigation, and regulatory monitoring.
- Independent from AI development teams to ensure objectivity in validation and review.
- Reports to the Chief Risk Officer with functional reporting to the Chief AI Officer.

10. Training Requirements

Meridian invests in comprehensive AI literacy and responsible AI training for all professionals:

- Responsible AI Foundations (mandatory, all staff): Annual 2-hour module covering AI ethics principles, acceptable use policy, data protection in AI, and incident reporting. FY2025 completion rate: 96.8%.
- GenAI for Professionals (mandatory, all client-facing staff): 4-hour module on effective and responsible use of generative AI tools in client engagements. Includes hands-on exercises with approved tools and case studies on AI risks.
- AI Risk Management (mandatory, AI Risk team and model developers): 16-hour certification program covering model risk management, validation methodology, fairness testing, and regulatory landscape.
- AI for Leaders (mandatory, Director level and above): 3-hour executive module on AI strategy, governance obligations, and leadership accountability for AI decisions.
- Specialized modules: Available for specific roles including AI/ML engineering, prompt engineering, AI in audit, AI in

legal services, and AI vendor management.

- Continuous learning: Monthly 'AI at Meridian' webinar series featuring internal practitioners, client case studies, and external guest speakers. Average attendance: 4,200 per session.